

Stop Being the Switchboard.

FREE GUIDE · OCTOBER 2026

Connect your AI agent to your production tools, so it plans, generates, edits and files the work while you direct it.

A practical setup guide for producers, editors, creators and studio teams.

Claude Code · Codex · fal · Runware · Higgsfield · Runway · Kling · ElevenLabs · DaVinci Resolve · Premiere · Final Cut · Blender · Descript · Canva · Figma · Gmail · Notion · Slack

Inspired Creative Group · inspiredcreativegroupinc.com
The Production Brain · theproductionbrain.ai

CONTENTS

What's inside.

Tap a line to jump to it. Start with "Where you type it" if you've never used a terminal.

How it works	3
Where you type it	4
Connecting a tool, click by click	5
1 · Install your agent	6
2 · Aggregators: fal, Runware	9
3 · Higgsfield, Runway, Kling	11
4 · ElevenLabs: voice and music	13
5 · Your editor: Resolve, Premiere, Final Cut	14
6 · Blender: 3D and previs	17
7 · Content tools: Descript, Canva, Figma, Adobe	18
The production office	20
What each model is good for	21
Plan before you prompt	24
The safety checklist	26
Your first hour	27

HOW IT WORKS

Your agent can hold the cables itself.

If you plan in one chat, paste prompts into a generation site, download, open another site for the voice and another for the music, then import every file into your editor, you are the switchboard: the person plugging every cable by hand.

Connected, your agent sends each job, files the results and keeps a record of what it did. You stay where your experience counts.

MCP: the plug

MCP (Model Context Protocol) lets your agent use a tool directly. Your agent (Claude Code or Codex) is the hub; each tool you connect becomes a set of actions it can call. Most vendors host their own MCP server with a browser sign-in, so connecting one is a single line.

Skills: your house rules

MCP gives the agent hands. **Skills** give it your rules: which model for which shot, how files are named, what a shot must pass before it goes out. A skill is a folder with a `SKILL.md` file in it. Some vendors publish official skills you install with one command.

BEFORE YOU START

Where you type it.

Every command box in this guide is labelled with where it goes. There are three places:

In Terminal
On a Mac, press Cmd + Space, type **Terminal** and press Return. On Windows, open **PowerShell**. Lines that start with `claude`, `codex`, `npm`, `npx` or `curl` go here.

Inside Claude Code
Type `claude` in Terminal to start it. Lines that start with a slash, like `/mcp`, go here, and so does everything you ask in plain English.

In the Claude app (no terminal)
Claude on desktop or claude.ai can connect most tools with a few clicks. Boxes marked **In the Claude app** give the clicks; the next page shows both ways step by step.

Rather not type any of it? This guide comes with an **agent edition**: a file written for your agent. Give it to Claude Code or Codex and say: "*Read this, and walk me through connecting my tools, one at a time. Ask me before anything that costs money.*" No agent yet? Open the **Code** tab in the Claude desktop app and drop the file in.

Sources: support.claude.com, "Get started with custom connectors using remote MCP" · code.claude.com/docs/en/mcp

BEFORE YOU START

Connecting a tool, click by click.

A. The tool is in Claude's directory

IN THE CLAUDE APP

1. Open **Customize**, then **Connectors**.
2. Click **+** next to Connectors to open the directory.
3. Search for the tool (for example "Canva") and click it.
4. Click **Connect**, then sign in to that tool's account and allow access.

B. You have a server address from this guide

IN THE CLAUDE APP

1. **Customize** → **Connectors**, then **+**.
2. Choose **Add custom connector**.
3. Type a name, paste the address (for example `https://mcp.runwayml.com/mcp`), click **Continue**, then **Add**.
4. Click **Connect** and sign in to that tool.

Check it worked: start a new chat and ask "What can you do with Canva?" On a Team or Enterprise plan, an Owner has to enable a connector before you can connect it. The Free plan allows one custom connector. Connectors added here usually show up in Claude Code too, when you sign in with the same Claude account.

Sources: support.claude.com, "Use the connectors directory" and "Get started with custom connectors using remote MCP"

STEP 1 · THE HUB

Install your agent.

Claude Code

Needs a paid Claude plan (Pro, Max, Team, Enterprise) or a Console account. The free plan doesn't include Claude Code.

IN TERMINAL

```
curl -fsSL https://claude.ai/install.sh | bash
claude --version
claude doctor
```

Other routes: `brew install --cask claude-code`, or in Windows PowerShell `irm https://claude.ai/install.ps1 | iex`. Type `claude` to start it, and `claude mcp list` to see your connections.

Source: code.claude.com/docs/en/setup

STEP 1 · THE HUB

OpenAI Codex CLI

Included with ChatGPT Plus, Pro, Business, Edu or Enterprise, or usable with an API key.

IN TERMINAL

```
npm install -g @openai/codex
codex          # choose "Sign in with ChatGPT"
```

Codex adds tools the same way:

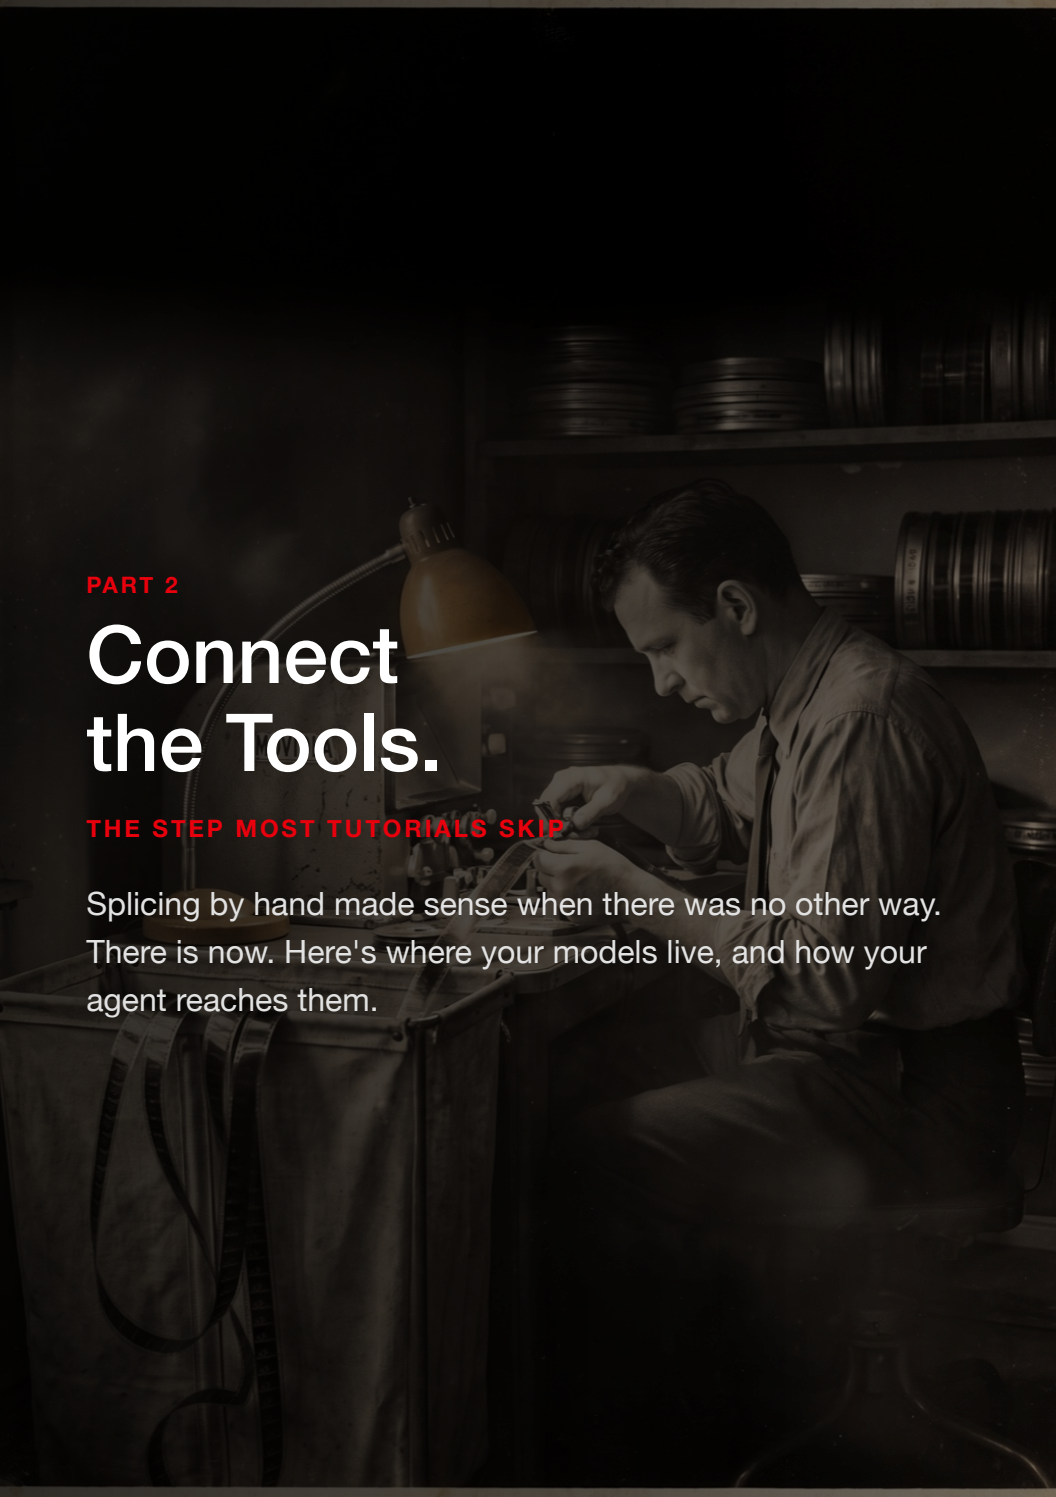
IN TERMINAL

```
codex mcp add <name> --url <server-url>
codex mcp login <name>
```

Settings live in `~/.codex/config.toml`.

Source: github.com/openai/codex · learn.chatgpt.com/docs/extend/mcp

Pick one hub to start. The rest of this guide uses Claude Code commands; every tool below also has a Codex equivalent.



PART 2

Connect the Tools.

THE STEP MOST TUTORIALS SKIP

Splicing by hand made sense when there was no other way. There is now. Here's where your models live, and how your agent reaches them.

STEP 2 · WHERE THE MODELS LIVE

Aggregators: many models, one connection.

An aggregator puts hundreds of image, video and audio models from different makers behind one connection and one bill. Your agent can compare and pick the right model for each shot instead of you opening a new site for each.

fal

A developer marketplace with a very large catalog (over a thousand models). Pay per use from prepaid credits. Each model's own terms apply to what you make with it.

IN TERMINAL

```
claude mcp add --transport http fal
https://mcp.fal.ai/mcp-relay
```

INSIDE CLAUDE CODE

```
/mcp
```

IN THE CLAUDE APP

Custom connector (way B) with the address `https://mcp.fal.ai/mcp-relay`.

Test: "Use fal to search for image generation models. Do not run a model."

Source: fal.ai/docs/model-apis/mcp

STEP 2 · WHERE THE MODELS LIVE

Runware

An inference platform with 400+ officially hosted models across image, video and audio (Veo, Kling, Seedance, Flux and ElevenLabs among them). Its terms are the clearest of the aggregators: you own your outputs, it doesn't train on them, and its partner models are covered for commercial use. It states SOC 2 and ISO 27001.

IN TERMINAL

```
claude mcp add --transport http runware
https://mcp.runware.ai
```

Then `/mcp` inside Claude Code. The sign-in page asks for your Runware API key once; the key stays on Runware's side.

IN THE CLAUDE APP

Custom connector (way B) with the address `https://mcp.runware.ai`.

Sources: runware.ai · runware.ai/terms

Pick	When
fal	You want the widest catalog and newest community models.
Runware	You do client work and want the cleanest commercial-use terms.

Watch out on video: the agent should submit a job, then check on it. Running the model again to "check progress" starts a new paid job.

STEP 3 · CREATIVE PLATFORMS

Higgsfield, Runway and Kling.

Creative platforms have their own apps and models. All three now let an outside agent generate with them, using the credits on your plan.

Higgsfield

Paid plan required. For Claude Code, Higgsfield recommends its own command-line tool and skills:

IN TERMINAL

```
npm i -g @higgsfield/cli
higgsfield auth login
npx skills add higgsfield-ai/skills
```

Or the hosted MCP:

IN TERMINAL

```
claude mcp add --transport http --scope user higgsfield
https://mcp.higgsfield.ai/mcp
```

IN THE CLAUDE APP

Custom connector (way B) with the address `https://mcp.higgsfield.ai/mcp`, then sign in with your Higgsfield account.

Know this: agent generations always use credits; "unlimited" applies only on higgsfield.ai itself.

Source: higgsfield.ai/creator-hub/help-center

STEP 3 · CREATIVE PLATFORMS

Runway

Hosted MCP since May 2026. Sign in with your Runway account; it generates image, video, music, speech and sound effects (Runway's models plus Veo, Kling, Seedance and others) and can run your saved Workflows and Brand Kits.

IN TERMINAL

```
claude mcp add --transport http --scope user runway
https://mcp.runwayml.com/mcp
```

IN THE CLAUDE APP

Custom connector (way B) with the address `https://mcp.runwayml.com/mcp`.

Know this: outside Enterprise, Runway's terms let it train on your inputs and outputs. Its developer API bills a separate wallet from your app plan.

Kling

Kling models straight from the source. Install Kling's command-line tool and skill, or its MCP (`https://kling.ai/mcp`), not both.

IN TERMINAL

```
npx skills add klingai-tech/skills
npm i -g @klingai/cli-global
kling login
```

Know this: only Personal workspace credits work through agents, result links expire after 24 hours, and the free plan allows no commercial use.

Sources: runway.com/mcp · runway.com/terms-of-use · kling.ai/document-api · kling.ai/docs/user-policy

STEP 4 · VOICE, EFFECTS AND MUSIC

ElevenLabs.

ElevenLabs

ElevenLabs runs a hosted MCP server. Sign in with your ElevenLabs account; no key needed.

> **IN TERMINAL**
>
> ```
> claude mcp add --transport http elevenlabs
> https://api.elevenlabs.io/v1/mcp
> ```

Then `/mcp` inside Claude Code to sign in. **Test:** "Generate a 10-second text-to-speech test line with a default voice."

> **IN THE CLAUDE APP**
>
> Directory (way A): search "ElevenLabs", Connect, sign in.

> If an older tutorial tells you to install the local `elevenlabs-mcp` server: it is deprecated. Use the hosted one above.

Music

Eleven Music composes a track from a prompt, from 3 seconds up to 10 minutes. The MCP docs don't list music tools yet, so music runs through the API, with your key in an environment variable. Several aggregators also carry music and speech models.

Sources: elevenlabs.io/docs (hosted MCP) · elevenlabs.io/docs/api-reference/music/compose

STEP 5 · THE EDIT

Your editor.

This changed fast. In September 2026, DaVinci Resolve added an official MCP server, built into the app.

DaVinci Resolve Studio

Native MCP server in **Resolve Studio 21.1** or later, made for Claude, Claude Code and Codex. It needs the paid Studio edition: from 21.1, the free version no longer allows outside scripting.

IN RESOLVE STUDIO

```
File → Setup AI Assistants
# restart Claude or Claude Code, then:
/mcp      (inside Claude Code)
```

What it can do: add markers and export them, rename clips, apply a LUT, reframe a timeline to 9:16, render deliverables in several formats.

Know this: Blackmagic says AI scripts get the same access to your disk and network as the assistant itself. Point it at a copy of the project first.

Sources: DaVinci Resolve 21.1 Reference Manual (Blackmagic Design) · digitalproduction.com, 8 Sept 2026

STEP 5 · THE EDIT

Premiere Pro: the route we use

Adobe has no official way for an agent to drive your Premiere timeline yet; its Claude connector only works on files you upload. The route we ran our own launch film's edit on is **Adobe Premiere Pro MCP**, a free community project. A panel inside Premiere takes the agent's instructions, so inside your open project it imports clips, builds sequences, places shots and adds markers, then reads it all back to check.

IN TERMINAL

```
git clone https://github.com/hetpatel-11/Adobe_Premiere_Pro_MCP ~/Adobe_Premiere_Pro_MCP
cd ~/Adobe_Premiere_Pro_MCP
npm run setup:mac
claude mcp add premiere -e PREMIERE_TEMP_DIR=/tmp/premiere-mcp-bridge -- node ~/Adobe_Premiere_Pro_MCP/dist/index.js
```

The setup line installs the panel and turns on the developer setting it needs (macOS). The last line connects it to Claude Code.

> **The point isn't a robot editor.** Let the agent do junior work: import, organize, lay a first pass from your notes. You start the edit at the part your experience is for.

STEP 5 · THE EDIT

IN PREMIERE

1. **Premiere Pro → Settings** (or Preferences) → **Plugins**: turn on **Enable developer mode**, then restart Premiere.
2. **Window → Extensions → MCP Bridge (CEP)**.
3. Set Temp Directory to `/tmp/premiere-mcp-bridge`, click **Save Configuration**, then **Start Bridge** and **Test Connection**.
4. Ask your agent: "List what you can do in Premiere." Then try it on a copy of a project.

Know this: it isn't made or supported by Adobe, and Premiere's scripting doesn't reach every menu, so expect assembly work, not a finished cut. Install it only from the original project page. Another free route, `adb-mcp`, does similar work through a newer plugin type.

No plugin? The agent can write an XML timeline that you import with **File → Import**. Slower, but nothing touches the app.

Final Cut Pro

A community server that works on exported FCPXML files: rough cuts, markers, trims, beat sync. You re-import the result.

IN TERMINAL

```
claude mcp add fcpxml -e FCP_PROJECTS_DIR=~/Movies -- uvx fcp-mcp-server
```

Sources: github.com/hetpatel-11/Adobe_Premiere_Pro_MCP · github.com/mikechambers/adb-mcp · helpx.adobe.com, "Adobe for Claude" · github.com/DareDev256/fcpxml-mcp-server

STEP 6 · 3D AND PREVIS

Blender.

For blockouts, previs, set extensions and camera moves you can hand to a video model as reference.

Blender (official)

Blender's own MCP server, from Blender Lab. Free, needs **Blender 5.1** or later.

> IN THE CLAUDE APP
>
> Directory (way A): search "Blender", Connect. Then install the Blender add-on the listing links to (drag it into Blender), and keep Blender open while you work.

Blender (community, more features)

The popular community server adds asset libraries (Poly Haven, Sketchfab) and AI 3D-model generators.

IN TERMINAL

```
claude mcp add blender uvx mcp-for-blender
uvx mcp-for-blender install-addon
```

Then enable the add-on in Blender. It sends anonymous usage data unless you set `DISABLE_TELEMETRY=true`.

> **Know this:** both run code the agent writes inside Blender, with no guards. Blender recommends a separate machine or a system with nothing sensitive on it. At minimum, work in a throwaway scene.

Sources: blender.org/lab/mcp-server · github.com/ahujasid/blender-mcp

STEP 7 · CONTENT TOOLS

Content tools.

For the cut-downs, captions, thumbnails and decks around every piece. Each is an official connector.

Descript
Edit by editing the transcript. Exports timelines to Resolve, Premiere, Final Cut and EDL, so it can feed your editor. **In the Claude app:** directory (way A), search "Descript".

IN TERMINAL

```
claude mcp add --transport http descript
https://api.descript.com/v2/mcp -s user
```

Canva
Design in your Canva account from the agent: social sizes, thumbnails, decks. **In the Claude app:** directory (way A), search "Canva", Connect, sign in.

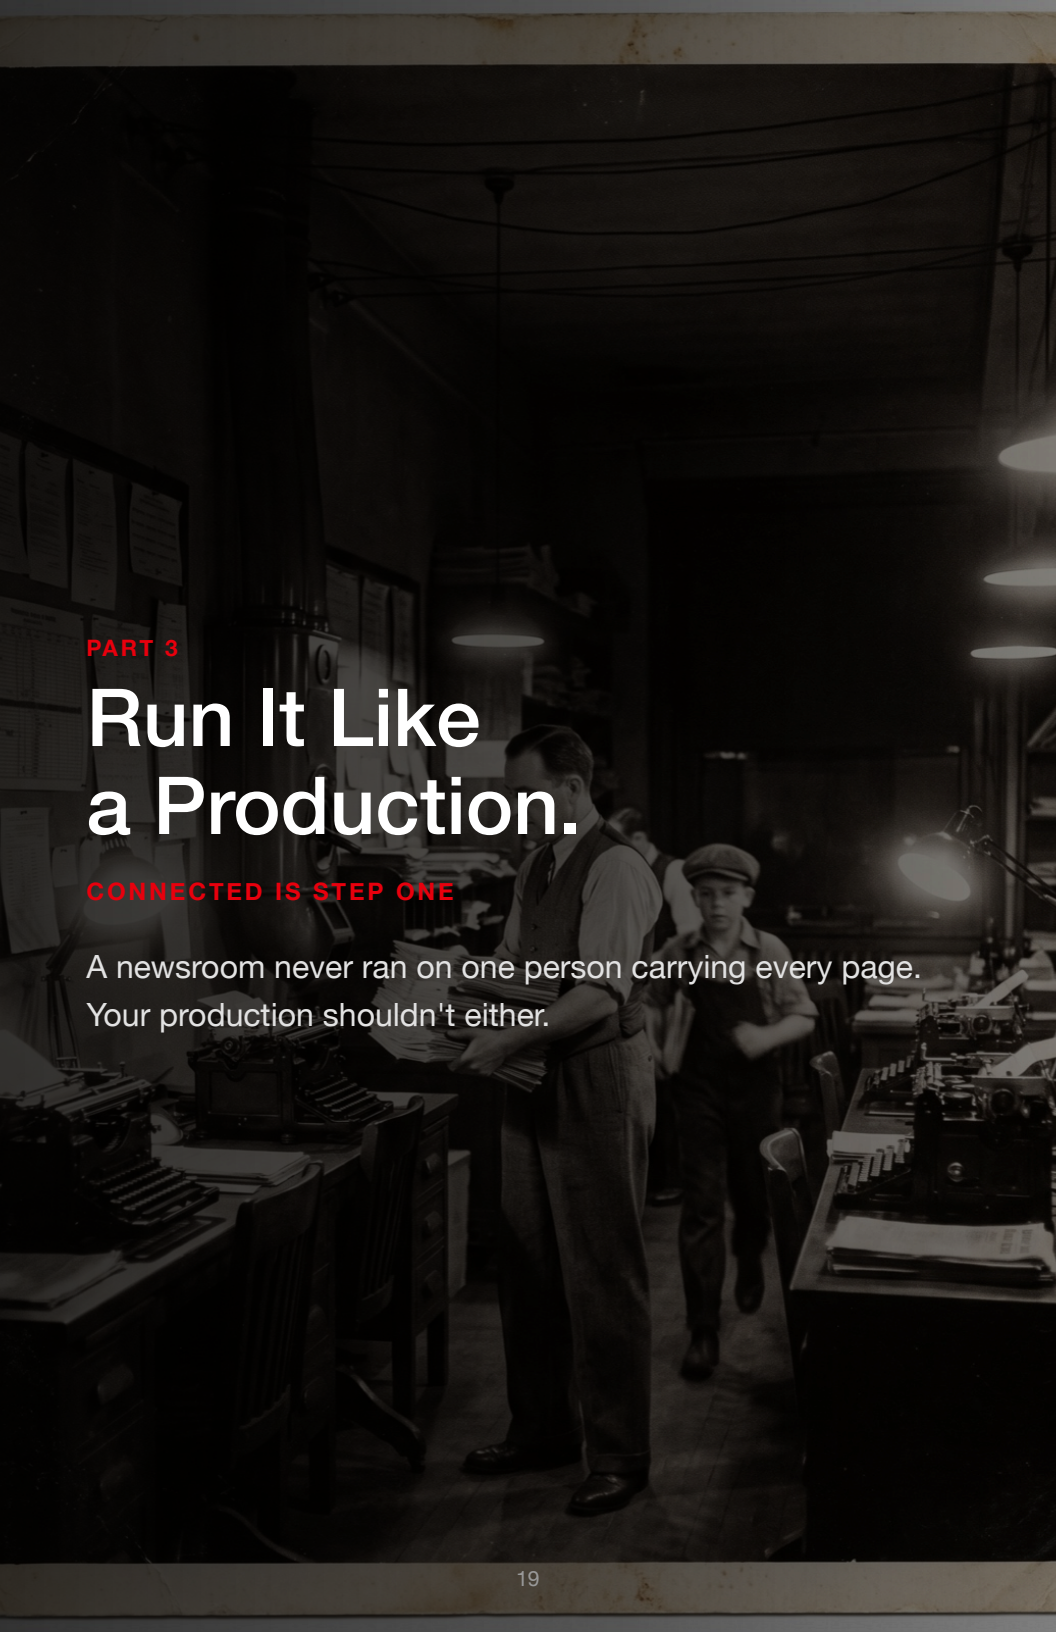
Figma
Storyboards, frames and design files.

IN TERMINAL

```
claude mcp add --transport http figma
https://mcp.figma.com/mcp
```

Adobe
One connector with Photoshop, Express and Firefly tools. Works on files you upload. **In the Claude app:** directory (way A), search "Adobe", Connect.

Sources: Descript, Figma and Adobe help centers · Claude connector directory



PART 3

Run It Like a Production.

CONNECTED IS STEP ONE

A newsroom never ran on one person carrying every page. Your production shouldn't either.

THE PRODUCTION OFFICE

Your agent runs the office too.

Most of a producer's day isn't generating. It's email, calendars, call notes and to-do lists. The same agent can run that side too.

All of these are in Claude's directory, so no Terminal is needed:

IN THE CLAUDE APP

Customize → **Connectors** → **+**, search the name (for example "Gmail"), click **Connect**, sign in with that account and allow access. Repeat for each one.

Job	Connector
Email, calendar, files	Gmail, Google Calendar, Google Drive
Tasks and docs	Notion
Team chat	Slack
Call notes	Granola, Otter or Fireflies

Try asking

- "Read yesterday's client call notes, list every change they asked for, and draft my reply. Don't send it."
- "Turn this shot list into Notion tasks with due dates."
- "Find 30 minutes next week for the colour review."

Drafts, not sends. Let the agent write; you press send. And never let it act on what an email tells it to do.

CHOOSING MODELS

What each model is good for.

Real examples, current October 2026. Models change monthly, so ask your agent what's newest before a big job.

Hero stills and style frames

TRY Nano Banana Pro (photoreal, up to 4K), Seedream 5.0 (up to 14 reference images), FLUX.2 (up to 10 references, exact brand hex colours).

AGENT CAN Run the same brief on all three and lay the results side by side.

CAN'T Tell you which one is on brand. You pick.

Text, logos and posters

TRY GPT Image 2 for copy in the image, in many languages. Recraft V4 Pro for logos as real vector (SVG) files.

AGENT CAN Spell the copy exactly as written and hand back a scalable logo.

CAN'T Get a transparent PNG out of GPT Image 2. Use Recraft for marks.

The same character across shots

TRY Seedance 2.0 and Kling 3.0 with reference images; Higgsfield's saved characters.

AGENT CAN Hold a face when you give it the same reference sheet every time.

CAN'T Guarantee wardrobe, hands and props across a long sequence. Check every shot.

Dialogue and sound in the shot

TRY Veo 3.1, Kling 3.0 and Wan generate speech and sound with the picture.

AGENT CAN Deliver a lip-synced line in one pass.

CAN'T Use your own brand voice. Record that in ElevenLabs and sync it in the edit.

Changing real footage

TRY Kling O3 video edit and Seedance edit modes (swap the environment, keep the move); Kling Motion Control and Higgsfield motion transfer (copy a performance onto a new character).

AGENT CAN Keep your camera move, light and timing while it replaces what's in frame.

CAN'T Hold fine detail through fast motion, or run long clips. Plan short pieces.

Long continuous takes

TRY Seedance 2.5: up to 30 seconds in one pass (1080p on Higgsfield, 720p on fal). Most models stop at 10 to 15 seconds.

AGENT CAN Generate one unbroken take instead of stitching clips.

CAN'T Make it cheap: it costs more per second than most models. Save it for the shot that needs it.

Hitting an exact end frame

TRY Veo 3.1 and Hailuo H3 first-and-last-frame: you give the start and end images, the model fills the move between them.

AGENT CAN Land a transition on the frame you designed.

CAN'T Invent a believable move between two unrelated frames.

CHOOSING MODELS

What the agent does, and what stays yours.

It can

- Compare models on price, limits and commercial terms before you spend.
- Write each model's prompt the way that model likes it.
- Submit jobs, track them and download results as they finish.
- File every result with its prompt, model, settings and cost.
- Run a bake-off of one shot on three or four models while you work on something else.

It can't

- Decide what's good. Taste, brand and the client's eye are yours.
- Push past a model's limits. A 15-second cap stays 15 seconds.
- Use website-only features, or the "unlimited" plans that only apply on a platform's own site.
- See its own results unless you give it a model that can look at frames, and even then you check.

THE WORKFLOW

Plan before you prompt.

Run every job in this order: **brief** (goal, audience, deliverables, deadline) → **plan** (shot list, which model per shot) → **generate** (agent checks the price first) → **file** → **assemble** (agent preps, you direct the cut).

Let the agent do the comparing

Ask it: "Find three models that can do this shot, show me their price per generation and their terms, then test the cheapest at low resolution." Keep the winner in a skill so the next production starts from what you learned.

Specialists and a validator

Give each model you use a **specialist** skill that knows how that model likes to be prompted. Then add a **validator**: a second agent, ideally on a different model, that checks every prompt against your rules before it goes out.

THE WORKFLOW

A folder your agent can follow

YOUR PROJECT FOLDER

```
my-project/
  00_brief/        brief.md
  01_plan/         shotlist.md, assets.md
  02_references/   images, sheets,
                   style frames
  03_generated/    one folder per shot;
                   each file saved with a
                   note: prompt, model,
                   settings, date, cost
  04_audio/        voice, effects, music
  05_edit/         timelines and exports
  06_delivery/     final files only
  decisions.md     what you chose, and why
```

Put this structure in your project's instructions file (`CLAUDE.md` for Claude Code, `AGENTS.md` for Codex) so the agent files everything without being asked.

BEFORE YOU CONNECT ANYTHING

The safety checklist.

- **Paid plans.** Claude Code and most generation tools need one for agent access.
- **Every call can cost.** Through an agent, generations use credits even where a website offers "unlimited". Ask the agent for the price first.
- **Test for free first.** A balance check or a model search before the first paid job.
- **Keys never go in chats or shared files.** Browser sign-in where offered, environment variables otherwise.
- **Read the terms for client work.** Who owns the output, whether commercial use needs a paid plan, whether the platform trains on your work.
- **Trust your servers.** Official or well-known servers only. A server that reads outside content (web pages, email) can carry instructions you didn't write.
- **Work on copies.** Let an agent touch a duplicate of a project or timeline first.
- **Download results right away.** Some platforms expire result links.

YOUR FIRST HOUR

Start small.

1. Install Claude Code or Codex (Step 1), or start in the Claude app with no terminal at all.
2. Connect one aggregator or platform and run its free test (Steps 2 and 3).
3. Connect ElevenLabs and make one test line (Step 4).
4. Create the project folder from the workflow page, then ask the agent to generate one shot and one voice line from a short brief and file both with their notes.
5. If you cut in Resolve Studio, connect it and ask the agent to add markers to a **copy** of a timeline.

That's the whole idea in miniature. Everything after is more of the same, planned.

Who Remembers What Worked?

THE PART CONNECTING DOESN'T SOLVE

Connected agents leave a trail: prompts, briefs, decisions, files. Months later, someone still has to know which prompt worked, which track was cleared, and what the client decided. Today that someone is you.

Inspired Creative Group built the Production Brain for exactly this. It's the centre of our own agentic production system, and it gives your studio the same thing: one memory on your own machine that every AI tool you use can read, checked by local models every night. Your studio's knowledge stays yours, whichever AI provider you use next.

> **One studio gets a custom build for free.** Applications close October 31: theproductionbrain.ai
>
> **Want your studio's workflow mapped?** Book a free 30-minute workflow review: message Juan Valencia on LinkedIn.

© 2026 Inspired Creative Group Inc. · inspiredcreativegroupinc.com
The Production Brain · theproductionbrain.ai
Setup steps checked against official documentation, 17 October 2026.